

# Travels with ICN

The road traversed and the road ahead

Dave Oran

Network Systems Research & Design / MIT



**NETWORK SYSTEMS RESEARCH & DESIGN**

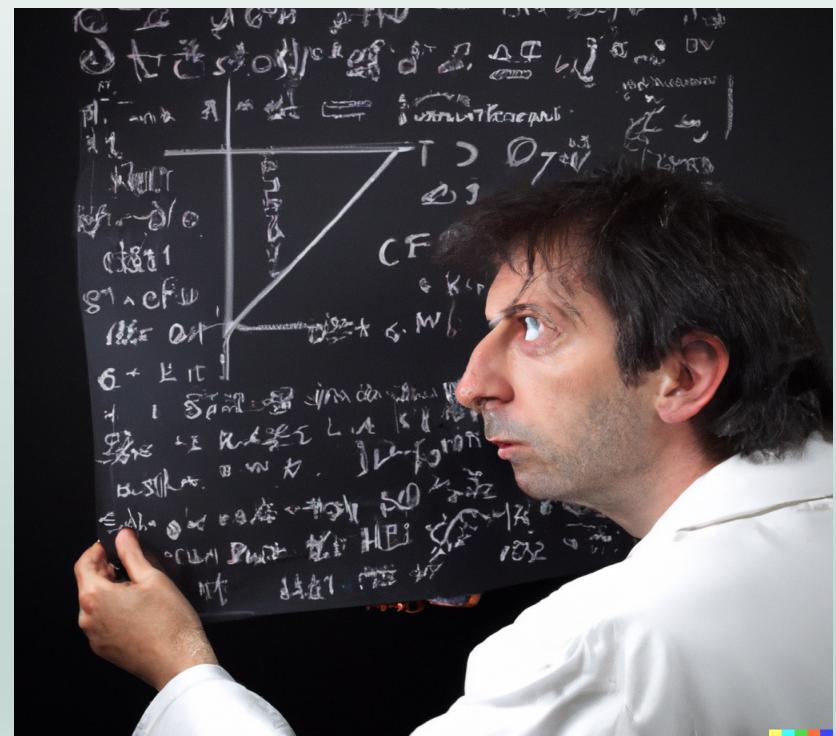
September 20, 2022  
ICN 2022, Osaka Japan

# Some Historical context – I've been around a while

Late 1960's	Satellite downlink data scrubbing	
Early 1970's	Banking Teller system + National Lottery (México)	
Late 1970's	DECnet + SNA Gateways	
Early 1980s	Internet Routing Protocols	
Late 1980s	RPC + Distributed Naming Services (OSF DCE)	
Early 1990s	Wireless LANs & Distributed File System	
1996-2006	VoIP	
2004-2010	IPTV	
2010-2016	Adaptive Internet Streaming	

# Why did I get interested in ICN (2011±)

- Potentially disruptive to the Networking business
- Opportunity to revisit hard problems in networking from different set of assumptions
- “Crazy” to keep per packet state in the network ...but... memory is a lot cheaper than it used to be



So, what did I learn (giving away the punchline)

Information-centric Networking

≠

Information-centric Networking

# Different Emphases

## Information-centric

- Naming forms (hierarchy, graph, attribute-based)
- Namespace design
- Named object integrity and confidentiality
- Trust schemata & Provenance
- Interaction of consumers & producers of data

## Networking

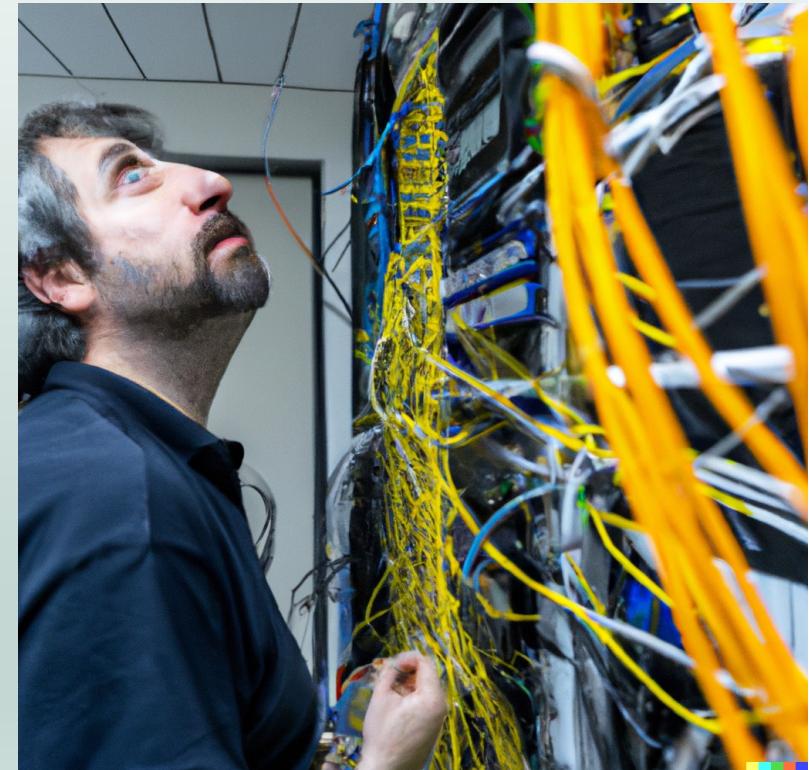
- Routing
- Forwarding
- Congestion Control
- Mobility
- Security of network devices
- DDoS prevention/mitigation
- Privacy against surveillance of network traffic

These are the things I'll talk about here

# Routing

# Name-based Routing is Seductive!

- Finally – location-independent routing using content names!
- Routers have lots more bits in the packet format to play with!
- But...
  - hierarchy doesn't help if names don't follow the network topology
  - ...and can we route on attributes? Or graph names?
  - Routing protocols scale as {State x Rate}



# We've tried lots of things for Name-based routing

- Link-State IGPs (e.g. NLSR)
- Multi-ring DHTs – (e.g. Liu, Foy, Zhang in ICN 2012)
- Bloom Filters – (e.g. Braun, Salamatian, Thomas in CCNC 2018)
- On-demand calculation – (e.g. Ascigil, Psaras, Pavlou) in ICN 2018)
- ...bunch of others
- But none of them have achieved:
  - Internet wide scale
  - Multi- AS operation
- So far, we've wound up punting to a Name Translation service
  - Fall back to DNS (sigh...)
  - Try a translation service that claims to scale (e.g, GNS or Mobility First)
  - Use Manifests containing multiple topological names

# Is there a lesson here?

**Lesson 1:** After 40+ years, no magic bullet to escape  
{State x Rate} scaling for routing

**Lesson 2:** All is not lost – lots of interesting  
applications don't need internet scale  
routing, but

**Lesson 3:** Few (any?) applications can get away with  
single AS administrative scope



# Forwarding

# The early approach to forwarding in CCNx & NDN

- Rich functionality was believed to be needed to demonstrate superiority of an L3 approach to ICN, for example:
  - Prefix matching against cached data allows content discovery to be “built in” to the base protocol
  - Selectors allow consumers to flexibly traverse collections
  - Exclusions needed to bypass data you really wish hadn’t wound up in the caches
- But don’t worry (yet) about performance, because...
- Doing premature optimizations to get performance risks freezing the architecture too soon

# I didn't believe this

- Making things go fast is interesting by itself
- You can't find "big O" problems if you don't care about performance
- Good to know what the bottlenecks are right away if you have very different resource tradeoffs (CPU, memory, bandwidth, latency) from conventional protocols



# So we took a real router and made ICN work

- Ran on Service Card of Cisco 9K Router
- We got ~1.3M PPS, and 20 Gbps
  - ...not bad for mostly software in 2013
- We learned a whole bunch of interesting things (next slides)
- A number of the key algorithms were adopted in later efforts (e.g. NIST forwarder)



# Some things we learned about forwarding, Part 1

- Router hardware was (is?) optimized for moving packets across busses/crossbars fast, not doing lots of computation or mutating state
- Memory bandwidth was the bottleneck, and seems to be the persistent bottleneck for this kind of forwarding across generations of hardware
- Algorithm selection the single biggest effect (yes, big O again!)
  - Hashing won over TRIES and other popular FIB approaches
  - Some fast hashes (e.g. CityHash64) were massively insecure, so chose a secure hash (SIPHash), which had side benefit of being restartable!
- Clever engineering also matters
  - Data structures that allow you to use SIMD instructions
  - Mix of long names and short names required some cleverness to bound worst case while having good average case forwarding lookup performance

# Some things we learned about forwarding, Part 2

- Three NDN and CCNx features were performance killers, so we gave up trying to implement them:
  - Prefix matching against CS blew memory access budget by 1-2 orders of magnitude
  - Ditto for Selectors
  - Exclusions (we didn't even try)

**Lesson 1:** Trying to go fast can inform protocol design even in early stages

**Lesson 2:** Remember that network routers/switches are not servers even if they have general purpose CPUs and DRAM to play with

# Congestion Control

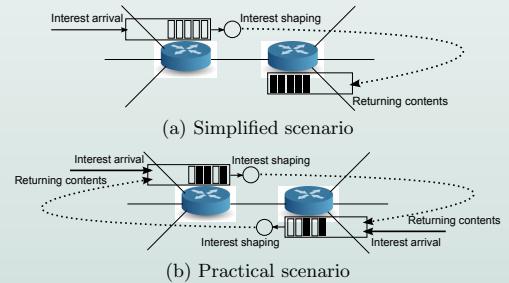
# Stateful forwarding makes Congestion Control *really* different from IP

- Hop-by-hop feedback
- Can get feedback at RTT to bottleneck rather than RTT end-to-end
- Not constrained by limited bits in IP header
- Symmetric routing means congestion state is also symmetric



# How do we exploit these properties to get better congestion control?

- We can shape traffic hop-by-hop
  - By shaping Interests we can control rate of returning Data
  - Can get much higher link utilization without loss than IP
- We can make rate-based schemes work!
  - Room in the packets for actual rate of bottleneck, cheap to compute and propagate back in Data
  - If congested, can NACK from bottleneck rather than dropping, giving shorter feedback RTT and retransmission timeliness
- We can police Interests when overloaded since we know rate of reverse link due to symmetric forwarding



An Improved Hop-by-hop Interest Shaper for Congestion Control in Named Data Networking

Yaogong Wang  
North Carolina State  
University, USA  
ywang5@ncsu.edu

Natalya Rozhnova  
Université Pierre et Marie  
Curie (UPMC), France  
natalya.rozhnova@lip6.fr

Ashok Narayanan  
Cisco Systems, USA  
ashokn@cisco.com

David Oran  
Cisco Systems, USA  
oran@cisco.com

Inpong Rhee  
Northeastern  
University, USA  
rhee@ncsu.edu

Jim Gibson  
Cisco Systems  
gibson@cisco.com

## MIRCC: Multipath-aware ICN Rate-based Congestion Control

Milad Mahdian  
Northeastern University  
mmahdian@ecu.edu

Somaya Arianfar  
Cisco Systems  
sarianfa@cisco.com

Dave Oran  
Cisco Systems  
oran@cisco.com

# Multipath rate based schemes practical

- We can learn multiple paths and maintain rate for each active path:
- We can steer packets explicitly onto paths
  - This allows proportional traffic splitting over paths

Path Switching  
in Content Centric and Named Data Networks

Ilya Moiseenko  
Cisco Systems  
ilmoisee@cisco.com

Dave Oran  
Network Systems Research & Design  
doran@random.net

# This rosy picture does have a few limitations

- From Interest arrival, you actually don't know how big a Data packet will come back, only that it will be at most one
  - Some ways to fix this, (e.g. consumer often knows and can say so in the Interest) but they haven't been tried
  - Ratio of Interest size to Data size affects efficiency – how much capacity for Interests on reverse link?
- No Magic for some universal congestion control complications:
  - End-to-end RTT uncertainty can be quite large – how pessimistic should we be? Long interest lifetimes make this even trickier.
  - If link rate changes faster than the link RTT (e.g. wireless) shaping errors will occur

# Not everything we tried worked out

- Congestion signaling hop-by-hop should make in-network retransmission attractive, but experiments seem to show it causes more problems than it solves
  - Inflates RTT artificially – consumer can't tell what part of RTT is propagation, accumulated queuing delay on successful path, versus delay via retransmission retries.
  - This also could be fixed, but at possibly significant complexity cost.

# Some lessons about ICN Congestion Control

**Lesson 1:** Congestion Control research is not a project, it's a career

**Lesson 2:** If you've paid the (high) price of per-packet state in the network, you should leverage it as much as possible

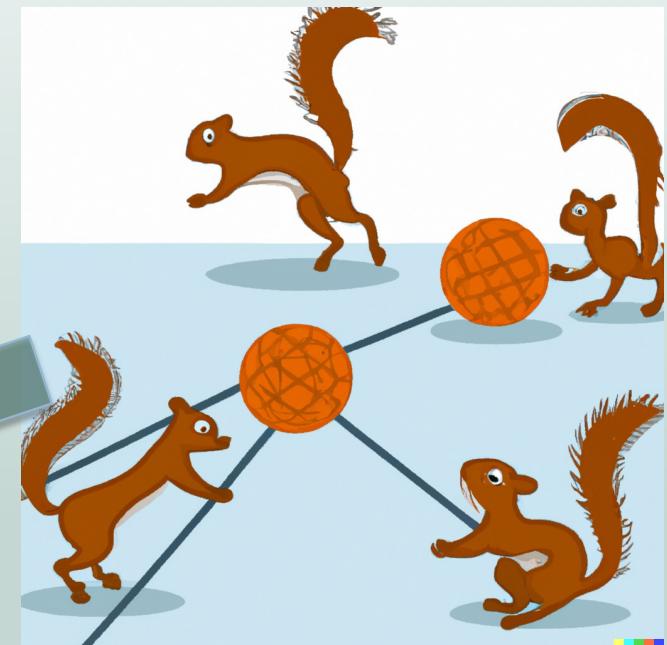
**Lesson 3:** None of this magically satisfies QoS fantasies, but it does protect the network from overload and can be made fair



# Interaction Models

# Interaction Models various applications use

- Content fetch
    - Been there, done that...
  - Synchronize shared data
    - Variety of Sync protocols developed for NDN  
not a "solved problem" yet
  - Publish/Subscribe
  - Remote Invocation (aka RPC) & Distributed Computations
    - Some useful early research through NEN and Nfaas
    - More recent work with RICE, CEN, and RF
  - Restful Transactions (aka Web)
- (This is what I'll talk about hereout)*
- I'll talk a bit about this later*



# Computations are different from fetching data

- Need to deal with different timescales
  - Computations can run for many network RTTs
  - Long Interest Lifetimes suck – too much state and not robust against network failures
- Input Data matters – do you Push it?
  - Too much unsolicited data screws up congestion control
- Do you trust the caller?
  - Proving identity via request sacrifices consumer anonymity
- Most interactions need more than a two-way handshake to avoid expensive retries to ascertain if something worked.



# Evolution of research into ICN for Distributed Computing: RICE, CFN, and RF

- RICE (ICN'18) addressed
  - Avoiding big request messages
  - Separating computation initiation from result delivery
  - Client authorization without losing anonymity to the network
  - Clean support of both idempotent and non-idempotent functions
- Compute First Networking (ICN'19) addressed
  - How to name computations and generate the interaction procedure directly from application code
  - How to orchestrate the placement of functions with Name-based routing
- Reflexive Forwarding (ICNRG Draft'22) cleans up and generalizes the handshake protocol machinery that RICE had hacked together
  - Lighter weight state management using “PIT Tokens”
  - Cleaner naming convention for the Interests going from producer to consumer

RICE: Remote Method Invocation in ICN			
Michał Król UCL m.krol@ucl.ac.uk	Karim Habak Georgia Tech karim.habak@gatech.edu	David Oran Network Systems Research & Design daveoran@random.net	
Dirk Kutscher Huawei dirk.kutscher@huawei.com	Ioannis Psaras UCL i.psaras@ucl.ac.uk		

Compute First Networking: Distributed Computing meets ICN			
Michał Król University College London/UCLouvain michal.krol@uclouvain.be	Spyridon Mastorakis University of Nebraska, Omaha smastorakis@unomaha.edu	David Oran Network Systems Research & Design daveoran@random.net	Dirk Kutscher University of Applied Sciences Emden/Leer dirk.kutscher@hs-emden-leer.de

Workgroup:	ICNRG
Internet-Draft:	<a href="#">draft-oran-icnrg-reflexive-forwarding-03</a>
Updates:	<a href="#">8569, 8609</a> (if approved)
Published:	6 June 2022
Intended Status:	Experimental
Expires:	8 December 2022
Authors:	D. Oran <i>Network Systems Research and Design</i> D. Kutscher <i>University of Applied Sciences Emden/Leer</i>

Reflexive Forwarding for CCNx and NDN Protocols

# Some Lessons learned from working on distributed computing with ICN

**Lesson 1:** Content fetch with two-way handshakes is a poor match for doing distributed computations.

**Lesson 2:** Extensions to the base protocols can give a flexible underpinning for multiple interaction models



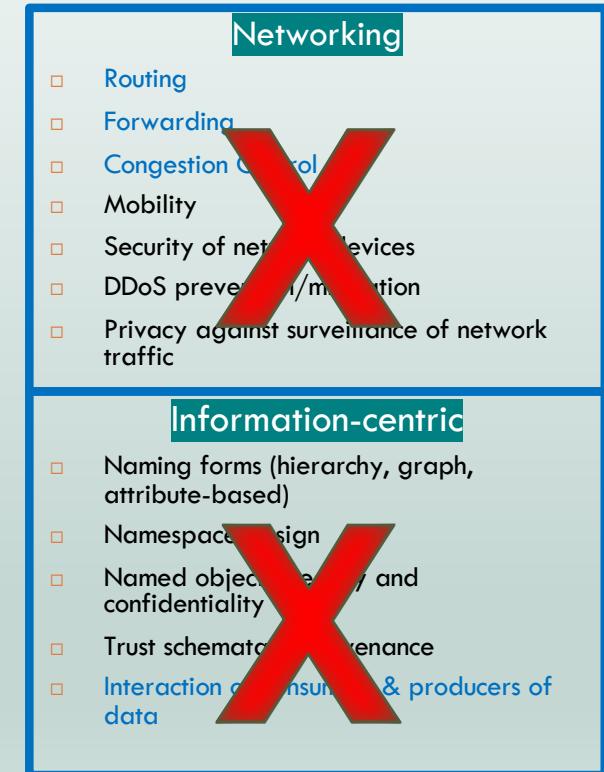
# We have two “extreme” data points in the design space to learn from

If you only care about the Network part, look at hICN:

- Use IPv6 packets and steal bits from transport for demux
  - No fancy naming
  - No object security built in
- Get most of the nice routing, forwarding & Congestion control properties of NDN and CCNx

If you only care about the Information-centric part, look at OSCORE:

- Hierarchical naming, Secure objects, Web-like semantics
- But no help from the network
  - Need proxies and L7 caches
  - Interactions tied to transactions rather than objects themselves



# Looking Ahead

What can we learn from our research to guide future research?

Are the lessons transferable?

**WARNING!**  
**Speculation Alert!!!**

# What belongs in which Layers?

**Explore the tradeoffs between doing things at the application layer versus pushing functionality down, perhaps to L2 as well as L3.**

- Experimentation with APIs seems central to answering this question
- As are choosing the right metrics: flexibility, throughput, latency, resilience



# Application design & Interaction models

**Can we have a “narrow waist” that works well for multiple Interaction models?**

- We know it works for content fetch
- What about:
  - Sync
  - RPC / Distributed Computing
  - Restful Web transactions
  - Pub/Sub
- Is the multi-way handshake machinery of Reflexive Forwarding adequate (or right) or do we need something else?



# Yes, Congestion Control still appeals

**Congestion control is a perennial area with many subtle issues to work on**

- How much to trust applications to respond to congestion signals (or packet drops) versus enforcement by network protocols
- Multipath congestion control is still ripe for improvements
- Dare I say it? QoS machinery for ICN.



# Computing in the Network?

**Should we instantiate transport functions and knowledge of interaction models in network devices (ala COIN)?**

- What ICN functions might belong in switches? What can be done in the face of the constrained programming model?
- Do we even need this? Why can't everything be done in servers?



# Is multi-destination delivery a winner in practice?

## IP Multicast hasn't been a raging success. Is the same true for multi-destination delivery in ICN?

- ICN multi-destination delivery seems a big win.  
Can we convincingly demonstrate that?
- On the other hand, Interest Multicast (for things like Sync interactions) seems to suffer from the same pathologies as IP multicast. Can this be overcome?



# Summing up the lessons

Information-centric networking (ICN) is an approach to networking that focuses on the distribution and management of information, rather than on the distribution of packets. ICN has been the subject of extensive research over the past ten years, and a number of lessons have been learned from this research.

One of the key benefits of ICN is that it can improve the efficiency of content delivery. In particular, ICN can reduce the amount of bandwidth that is required for content delivery, and can also improve the reliability of content delivery.

ICN can also improve the security of content delivery. In particular, ICN can help to protect against attacks that exploit vulnerabilities in the network infrastructure. ICN can also help to improve the privacy of content delivery. In particular, ICN can help to protect the privacy of users by preventing third-party observers from tracking the activities of users.

ICN can also help to improve the scalability of content delivery. In particular, ICN can help to improve the scalability of content delivery by allowing content to be cached closer to the users who need it.

**OOPS...this was  
generated by GPT-3  
DENIED**

# Let's try again...

- How can the networking insights we've gained from ICN protocols inform the construction of Information Centric systems and applications?
  - Whether and how to utilize name-based routing to achieve robustness and performance scaling for distributed applications?
  - Where does caching help or not help and how to best utilize caches?
  - Does pushing Names down to lower layers help latency? Resilience? Fairness?
- How can the insights we've gained from applying Information Centricity in applications inform what we bother to change the network to do, and what not?
  - Do things like multipath forwarding, in-network retransmission, or reflexive forwarding actually enable applications that are hard or infeasible to do without them?
  - Is there a big win for wireless networks in terms of optimizing a scarce resource or having more robust and responsive mobility characteristics?

## Revisiting the Punchline

ICN is  
Information-centric  $\cup$  Networking  
not  
Information-centric  $\cap$  Networking

# None of this would be possible without:

Alex Afanasyev	Cesar Ghali	Luca Muscariello	Ersin Uzun
Somaya Arianfar	Jim Gibson	Ashok Narayanan	Yaogong Wang
Mark Baugher	Cenk Gündoğan	Börje Ohlman	Matthias Waehlisch
Hila Ben Abraham	Karim Habak	Christos Papadopoulos	Cedric Westphal
Jeff Burke	Michał Król	Ioannis Psaras	Bastiaan Wissingsh
Ken Calvert	Dirk Kutscher	K.K. Ramakrishnan	Chris Wood
Giovanna Carofiglio	Yuanjie Li	Natalya Rozhnova	Edmund Yeh
Taejoong Chung	Milad Mahdian	Thomas Schmidt	Haowei Yuan
Alberto Compagno	Spyros Mastorakis	Eve Schooler	Lixia Zhang
Patrick Crowley	Maziar Mirzazad	Wentao Shang	
Ralph Droms	Satyajayant Misra	Won So	
Serge Fdida	Ilya Moiseenko	Mark Stapp	(with Apologies to
J.J. Garcia Luna Aceves	Marie-Jose Montpetit	Christian Tschudin	those I'm sure I
Mevlut Garip	Marc Mosko	Gene Tsudik	missed!)

# Thanks for Listening!!\*

Many thanks to Dirk Kutscher & Ken Calvert for sage advice on the important things to get across, and many great suggestions

\*Most photos are *not* real – they were generated by DALL-E 2 - can you guess which?